

PnS 2018

Deep Learning with Raspberry Pi



Session 3

PnS 2018 Team

Institute of Neuroinformatics
University of Zürich and ETH Zürich



Outline

Artificial Neuron: Overview




- A basic computational model of the biological model
- Single neuron as linear/logistic regression

Multi-Layer Perceptron



- Neurons in an acyclic feed-forward graph
- Fully connected layers
- Each fully connected layer computation is a matrix multiplication, matrix addition and an activation function

What can an MLP learn?






- Neural Networks with at least one hidden layer are universal approximators¹
- More neurons are expected to approximate better

¹Approximation by superpositions of a sigmoidal function, by Cybenko G.
<http://cs231n.github.io/neural-networks-1/>

Regularization

- Overfitting more probable with larger models
- Could be prevented by using a regularization term in the loss function

Regularization

 $\lambda = 0.001$  $\lambda = 0.01$  $\lambda = 0.1$ 

- Use bigger networks but take measures to prevent overfitting

Working with images

- MLPs do not work well with images
- Hierarchy of local spatial features
- Extract these local spatial features through filters